

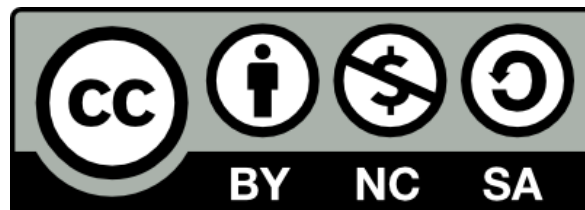
Software Install and Containers on ARCHER

Adrian Jackson

a.jackson@epcc.ed.ac.uk



Reusing this material



This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License.

<https://creativecommons.org/licenses/by-nc-sa/4.0/>

This means you are free to copy and redistribute the material and adapt and build on the material under the following terms: You must give appropriate credit, provide a link to the license and indicate if changes were made. If you adapt or build on the material you must distribute your work under the same license as the original.

Note that this presentation contains images owned by others. Please seek their permission before reusing these images.

Partners



Engineering and
Physical Sciences
Research Council

Natural
Environment
Research Council



THE UNIVERSITY
of EDINBURGH



**Hewlett Packard
Enterprise**

Installing software on ARCHER2

- ARCHER2 has a lot of software already installed
 - Using `module` system
 - `module avail`
 - `module list`
 - `module load PrgEnv-gnu`
 - Optimised central installs are preferable
 - `module load cray-python`
 - Currently version 3.8.5.0
 - Contains numpy, scipy, mpi4py, dask
 - `module load cray-R`
 - Currently version 4.0.3.0
- Two challenges to installing software
 - Cray Linux may be non-standard, and may not have dependencies installed
 - Default `home` directory installs won't work for the compute nodes

Scientific libraries

- Cray libsci loaded by default:
`module list`
`module show cray-libsci`
- Provides optimised:
 - blas
 - lapack
 - scalapack
- More info available on the system:
`man intro_libsci`
`man intro_blas1` `intro_blas2` `intro_blas3`
`man intro_lapack`
`man intro_scalapack`
- Other scientific software packages pre-installed
`module avail`
 - FFTW, HDF5, NetCDF, ADIOS
 - ARPACK, Boost, Eigen, GLM, HYPRE, Matio, Intel MKL, MUMPS, PETSc, SLEPc, Trilinos, SuperLU/SuperLU_DIST
 - Metis/Parmetis, Scotch/PT-Scotch

Compilers

- Three different compiler environments available on ARCHER2:
 - AMD Compiler Collection (AOCC)
`module load PrgEnv-aocc`
 - Based on clang and flang
 - GNU Compiler Collection (GCC)
`module load PrgEnv-gnu`
 - HPE Cray Compiler Collection (CCE) (default)
`module load PrgEnv-cray`
 - Cray Fortran compiler and clang for C/C++
- Compilation undertaken using:
 - `cc`, `CC`, `ftn`
- Different compiler versions are also available
 - i.e. `module swap gcc gcc/11.2.0`
`module avail cce`
`module avail gcc`
`module avail aocc`

Installing software

- Installing your own software for use on the compute nodes
 - Remember only `/work` is available on the compute nodes
- Python

```
module load cray-python
export PYTHONUSERBASE=/work/t01/t01/auser/.local
```

 - You will need to change `t01` to the project code for your project, and `auser` to your username

```
export PATH=$PYTHONUSERBASE/bin:$PATH
```

 - Can use virtual environments, i.e.:

```
source <<path to virtual environment>>/bin/activate
```
 - Can then use pip or conda to install software

```
pip install --user dask distributed
```
 - `--user` isn't required if virtual environments are being used
- Some python installs may need flags set for compilation

```
export CC=cc
export CXX=CC
export FC=ftn
```
- BLAS, LAPACK libraries provided from the `cray-libsci` module (loaded by default)

Installing software

- R

```
export R_LIBS_USER=/work/z19/z19/adrianj/Rinstall
R
install.packages('snow')
```
- May also need to configure install environment:
 - Create a preference directory

```
mkdir ~/.R
```
 - Add this to:

```
~/.R/Makevars
```
 - With the following lines:

```
CC = cc
CXX = CC
FC = ftn
```
 - Then can install from the command line:

```
R CMD INSTALL Rmpi_0.6-9.2.tar.gz --configure-args=" --with-Rmpi-
type=CRAY"
```


- Packages in library '/opt/R/4.0.3.0/lib64/R/library':
 - base The R Base Package
 - boot Bootstrap Functions (Originally by Angelo Canty for S)
 - class Functions for Classification
 - cluster "Finding Groups in Data": Cluster Analysis Extended Rousseeuw et al.
 - codetools Code Analysis Tools for R
 - compiler The R Compiler Package
 - datasets The R Datasets Package
 - foreign Read Data Stored by 'Minitab', 'S', 'SAS', 'SPSS', 'Stata', 'Systat', 'Weka', 'dBase', ...
 - graphics The R Graphics Package
 - grDevices The R Graphics Devices and Support for Colours and Fonts
 - grid The Grid Graphics Package
 - KernSmooth Functions for Kernel Smoothing Supporting Wand & Jones (1995)
 - lattice Trellis Graphics for R
 - MASS Support Functions and Datasets for Venables and Ripley's MASS
 - Matrix Sparse and Dense Matrix Classes and Methods
 - methods Formal Methods and Classes
 - mgcv Mixed GAM Computation Vehicle with Automatic Smoothness Estimation
 - nlme Linear and Nonlinear Mixed Effects Models
 - nnet Feed-Forward Neural Networks and Multinomial Log-Linear Models
 - parallel Support for Parallel computation in R
 - rpart Recursive Partitioning and Regression Trees
 - spatial Functions for Kriging and Point Pattern Analysis
 - splines Regression Spline Functions and Classes
 - stats The R Stats Package
 - stats4 Statistical Functions using S4 Classes
 - survival Survival Analysis
 - tcltk Tcl/Tk Interface
 - tools Tools for Package Development
 - utils The R Utils Package

Rscript -e "installed.packages()"

Containers

- Containers allow separation of kernel and user space for operating systems and applications
 - Enable different user space configurations for a given kernel space
 - Interface level between the two
 - Virtualises the hardware and operating system from the user software perspective
 - Lighter weight than full virtualisations (VM) but less isolated
- Docker is an example
 - Docker images are the files and directories that a docker container will be created from
 - Container is a runtime image (lightweight virtual machine)
 - Docker images can be obtained from the Docker hub
<https://hub.docker.com/>

Singularity

- Docker has some security configuration issues that restrict its use on shared resources like ARCHER2
- HPC systems also have specific requirements around optimised filesystems and networks
 - Singularity is a container implementation designed for HPC systems
 - Others are also available (shifter, Charliecloud, etc....)
- Singularity has two different versions
 - <https://apptainer.org/>
 - <https://sylabs.io/>
- Singularity on ARCHER2
module load singularity
singularity --version

Using singularity

- Download and run a docker container

```
singularity pull python-3.9.9.sif docker://python:3.9.9-slim-buster
singularity run python-3.9.9.sif python -c "print('hello')"
```
- Can also get containers from <http://datasets.datalad.org/?dir=/shub>

```
singularity pull hello-world.sif shub://vsoch/hello-world
singularity run hello-world.sif
```
- Different ways to run things using singularity

```
singularity run hello-world.sif
singularity shell hello-world.sif
singularity exec python-3.9.9.sif python -c "print('hello')"
```
- Check the default command/execution

```
singularity inspect -r hello-world.sif
```
- Some containers also available on ARCHER2

```
/work/y07/shared/singularity-images
```

Singularity users and files

- By default singularity (on ARCHER2) will bring in your user and groups from the host system

```
adrianj@ln04:~> singularity shell -B /work/z19/z19/adrianj hello-world.sif
Singularity> whoami
adrianj
Singularity> groups
z19 archer2-tds-login archer2-4c-login castep-admin cse-admin archer2-login
• Enables file access and permissions etc... to be maintained
```

- Filesystem is not imported...

```
• ...except your home directory
adrianj@ln04:~> singularity shell hello-world.sif
Singularity> pwd
/home/z19/z19/adrianj
```

- Can bring in additional directories using `-B` flag

```
adrianj@ln04:~> singularity shell -B /work/z19/z19/adrianj hello-world.sif
• Can specify bind path:
adrianj@ln04:~> singularity shell -B /work/z19/z19/adrianj:/workdir hello-world.sif
```

- Container is read only except your home directory and bind directories

Parallel applications using Singularity

- Node local (shared memory) applications
 - Singularity usage model the same as normal applications

```
#!/bin/bash --login
#SBATCH --job-name=my_app
#SBATCH --nodes=1
#SBATCH --tasks-per-node=1
#SBATCH --cpus-per-task=8
#SBATCH --time=00:10:00

#SBATCH --account=[budget code]
#SBATCH --partition=standard
#SBATCH --qos=standard
export OMP_NUM_THREADS=8
singularity run $SLURM_SUBMIT_DIR/my_app.sif
```
- Distributed memory applications (MPI) requires more care
 - Need MPI compatibility between host and container

Parallel applications using Singularity

```
#!/bin/bash
#SBATCH --job-name=singularity_parallel
#SBATCH --time=0:10:0
#SBATCH --nodes=2
#SBATCH --tasks-per-node=128
#SBATCH --cpus-per-task=1
#SBATCH --partition=standard
#SBATCH --qos=standard
#SBATCH --account=[budget code]
# Set the number of threads to 1.
# This prevents any threaded system libraries from automatically using threading.
export OMP_NUM_THREADS=1
# Set the LD_LIBRARY_PATH environment variable within the Singularity container
# to ensure that it used the correct MPI libraries.
export SINGULARITYENV_LD_LIBRARY_PATH= \
    /opt/cray/pe/mpich/8.1.9/ofc/gnu/9.1/lib-abi-mpich: \
    /opt/cray/pe/pmi/6.0.13/lib: \
    /opt/cray/libfabric/1.11.0.4.71/lib64: \
    /usr/lib64/host: \
    /usr/lib/x86_64-linux-gnu/libibverbs: \
    /.singularity.d/libs
# Set the options for the Singularity executable.
# This makes sure Cray Slingshot interconnect libraries are available
# from inside the container.
BIND_OPTS="-B /opt/cray,/usr/lib64:/usr/lib64/host,/usr/lib64/tcl"
BIND_OPTS="${BIND_OPTS},/var/spool/slurmd/mpi_cray_shasta"
# Launch the parallel job.
srun --hint=nomultithread --distribution=block:block \
    singularity run ${BIND_OPTS} osu_benchmarks.sif \
        collective/osu_allreduce
```

Parallel applications using Singularity

- Interactive compile

```
singularity run -B /work/z19/z19/adrianj:/workdir  
/work/y07/shared/singularity-images/mpich_base.sif mpicc -fopenmp  
-o /workdir/mpi_hello_world /workdir/mpi_hello_world.c
```

- Interactive run

```
srun --hint=nomultithread --distribution=block:block --nodes=1 --  
tasks-per-node=16 --cpus-per-task=8 --exclusive --  
partition=standard --qos=short --reservation=shortqos --  
account=z19 --time=0:20:0 singularity run "-B  
/work/z19/z19/adrianj:/workdir,/work/y07/shared,/opt:/opt,/usr/lib  
64:/usr/lib64/host,/usr/lib64/tcl,/var/spool/slurmd/mpi_cray_shast  
a" /work/y07/shared/singularity-images/mpich_base.sif  
/workdir/mpi_hello_world
```


MPI in Singularity

- Different modes for using MPI inside a Singularity container
 - <https://apptainer.org/user-docs/3.7/mpi.html#singularity-and-mpi-applications>
 - Host mode
 - Use the host MPI to run Singularity
 - Enables integration with batch system
 - Needs compatible MPI between host and container
 - Needs container to be configured to use high performance network
 - Bind mode
 - No MPI required within the container
 - Package application into the container
 - Specify where the host MPI is installed so can be accessed within the container

MPI in Singularity

- Host mode example:

- Build definition

Bootstrap: docker

From: ubuntu:20.04

%files

/home/singularity/osu-micro-benchmarks-5.8.tgz /root/

/home/singularity/mpich-3.4.3.tar.gz /root/

%environment

export SINGULARITY_MPICH_DIR=/usr

export OSU_DIR=/usr/local/osu/libexec/osu-micro-benchmarks/mpi

%post

apt-get -y update && DEBIAN_FRONTEND=noninteractive apt-get -y install build-essential libfabric-dev libibverbs-dev gfortran

cd /root

tar zxvf mpich-3.4.3.tar.gz && cd mpich-3.4.3

echo "Configuring and building MPICH..."

./configure --prefix=/usr --with-device=ch4:ofi && make -j8 && make install

cd /root

tar zxvf osu-micro-benchmarks-5.8.tgz

cd osu-micro-benchmarks-5.8/

echo "Configuring and building OSU Micro-Benchmarks..."

./configure --prefix=/usr/local/osu CC=/usr/bin/mpicc CXX=/usr/bin/mpicxx

make -j6 && make install

%runscript

echo "Rank \${SLURM_PROCID} - About to run: \${OSU_DIR}/\${*}"

exec \${OSU_DIR}/\${*}

- Build command

singularity build osu_benchmarks.sif osu_benchmarks.def

Creating images

- To modify images/build new images need
 - Root permissions on a system with singularity installed
 - Docker installed on a system (using a docker singularity container)

- Create image definition file

```
Bootstrap: docker
```

```
From: ubuntu:20.04
```

```
%post
```

```
apt-get -y update && apt-get install -y python3
```

```
%runscript
```

```
python3 -c 'import sys; print("Hello World! Hello from Python %s.%s.%s in our custom Singularity image!" % sys.version_info[:3])'
```

- Build image

```
singularity build my_test_image.sif my_test_image.def
```

- Or

```
docker run --privileged --rm --mount type=bind,source=${PWD},target=/home/singularity quay.io/singularity/singularity:v3.7.3-slim build /home/singularity/my_test_image.sif /home/singularity/my_test_image.def
```

- Can make more complex/functional images

- Different sections for definition files:
 - %setup, %files, %environment, %startscript, %test, %labels, %help
 - https://apptainer.org/user-docs/3.7/definition_files.html#sections
- Can sign containers for distributed etc...

Summary

- Plenty of software already available on ARCHER2
- Plus a range of compilers
- Installing your own R and Python libraries is straightforward
 - But ensuring they're available on the compute nodes requires configuration
- More complex software installs may benefit from containers
 - Basic container functionality simple
 - Interfacing with MPI and the /work filesystems require more care
- Ensuring software is as efficient as possible is important if using large amounts of compute time
 - Placement and binding for threads/processes important
 - Optimised maths libraries
 - New compilers and optimised compile options
 - etc...