

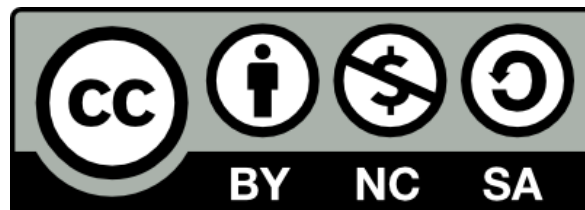
Parallel R and ARCHER2

Adrian Jackson

a.jackson@epcc.ed.ac.uk



Reusing this material



This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License.

<https://creativecommons.org/licenses/by-nc-sa/4.0/>

This means you are free to copy and redistribute the material and adapt and build on the material under the following terms: You must give appropriate credit, provide a link to the license and indicate if changes were made. If you adapt or build on the material you must distribute your work under the same license as the original.

Note that this presentation contains images owned by others. Please seek their permission before reusing these images.

Partners



Engineering and
Physical Sciences
Research Council

Natural
Environment
Research Council



THE UNIVERSITY
of EDINBURGH



Hewlett Packard
Enterprise

Parallel R

- Different motivations for parallelising R
 - Too much computational work
 - My program takes too long
 - Too much data to process
 - My program is quick enough, but I've too much data to process
 - Data is too large
 - My data set is too big to fit into memory
- Wide range of options:
 - <https://cran.r-project.org/web/views/HighPerformanceComputing.html>
 - Implicit parallelism
 - Explicit parallelism
 - Applications
 - Batch system integration, hardware resources, distributed computing integration
 - Native code interfaces
 - Big memory

Parallel options for R

- Shared memory parallelisation
 - Trivially parallel
 - (l)apply
 - Map
 - Loop parallelism
 - foreach
 - Threads
 - Rdsm
- Distributed memory parallelisation
 - parallel
 - snow
 - Rmpi
 - pbdR
 - foreach

Automatic parallelism

- R has some parallelism built in
- BLAS libraries and other underlying functionality already using shared memory parallelism, i.e.:

```
A <- matrix( rnorm(n*n), ncol=n, nrow=n )
B <- matrix( rnorm(n*n), ncol=n, nrow=n )
C <- A %*% B
```

```
> library("data.table")
```

```
data.table 1.14.2 using 128 threads (see ?getDTthreads). Latest news: r-  
datatable.com
```

- Can control number of workers on ARCHER2 using

```
export OMP_NUM_THREADS=
```

i.e.

```
export OMP_NUM_THREADS=2
```

...

```
> library("data.table")
```

```
data.table 1.14.2 using 2 threads (see ?getDTthreads). Latest news: r-  
datatable.com
```

- By default will use all available workers
 - Can be the wrong thing to do

Trivial trivial parallelism

- Task farms
 - Same code executed independently
 - Different initial conditions or different data sets
 - Slurm array jobs can support this

```
args <- commandArgs(TRUE)
set.seed(args[1])
data <- read.csv('dataset.csv')
result <- kmeans(data, centers=10, nstart=100)
print(result)
```

```
Rscript kmeans-trivial.R 1 > 1_output.dat
```

- Requires manual processing of the data
- Requires some care if using random number generators

Trivial parallelism

- Split, Apply, Combine
 - Partition data into groups to be worked on
 - Apply operations to each group
 - Combine results with some form of reduction at the end
- Requires making some workflows more complex to enable parallelisation:

```
data <- read.csv('dataset.csv')
result <- kmeans(data, centers=4, nstart=100)
print(result)
```

```
data <- read.csv('dataset.csv')
parallel.function <- function(i) {
  kmeans( data, centers=4, nstart=i )
}
results <- lapply( c(25, 25, 25, 25), FUN=parallel.function )

temp.vector <- sapply( results, function(result) { result$tot.withinss } )
result <- results[[which.min(temp.vector)]]
```


Combining results

- Some of the functionality we can use can create lists of answers, i.e.:

```
> res <- lapply(1:3, function(i) {  
+   sqrt(i)*sqrt(i*2)  
+ })
```

```
>
```

```
> print(res)
```

```
[[1]]
```

```
[1] 1.414214
```

```
[[2]]
```

```
[1] 2.828427
```

```
[[3]]
```

```
[1] 4.242641
```

- Can combine the results to produce a single vector or results, i.e.:

```
> do.call('c', res)
```

```
[1] 1.414214 2.828427 4.242641
```

mcapply

```
library(parallel)
data <- read.csv('dataset.csv')
parallel.function <- function(i) {
  kmeans( data, centers=4, nstart=i )
}
results <- mclapply( c(25, 25, 25, 25), FUN=parallel.function )

temp.vector <- sapply( results, function(result) { result$tot.withinss } )
result <- results[[which.min(temp.vector)]]
print(result)
```

- Uses processes rather than threads on ARCHER2.
- Uses **MC_CORES** environment variable to control parallelism (not **OMP_NUM_THREADS**)
 - Default to 2, can change, i.e.:
`export MC_CORES=16`
- Uses **mc_cores** argument as well:
`results <- mclapply(c(25, 25, 25, 25), FUN=parallel.function, mc_cores=16)`
- Can modify task scheduling/load balancing etc... with further arguments:
`results <- mclapply(c(25, 25, 25, 25), FUN=parallel.function, mc_cores=16,
mc.preschedule=FALSE)`

mcparallel mcollect

```
library(parallel)
p <- mcparallel(1:10)
q <- mcparallel(1:20)
# wait for both jobs to finish and collect all results
res <- mcollect(list(p, q))
```

```
library(parallel, quiet=TRUE)
source("data/airline/read_airline.R")
jan2010 <- read.airline("data/airline/airOT201001.csv.gz")
unique.planes <- mcparallel( length( unique( sort(jan2010$TAIL_NUM) ) ) )
median.elapsed <- mcparallel( median( jan2010$ACTUAL_ELAPSED_TIME, na.rm=TRUE ) )
ans <- mcollect( list(unique.planes, median.elapsed) )
ans
```

foreach

```
library(foreach)
data <- read.csv('dataset.csv')
results <- foreach( i = c(25,25,25,25) ) %do% {
  kmeans( x=data, centers=4, nstart=i )
}
temp.vector <- sapply( results, function(result)
  { result$tot.withinss } )
result <- results[[which.min(temp.vector)]]

print(result)

...
x <- foreach(i=1:4, .combine='cbind') %do%
rnorm(4)
```

foreach parallel

```
library(foreach)
library(doMC)
data <- read.csv('dataset.csv')

registerDoMC(4)
results <- foreach( i = c(25,25,25,25) ) %dopar% {
  kmeans( x=data, centers=4, nstart=i )
}

temp.vector <- sapply( results, function(result)
  { result$tot.withinss } )
result <- results[[which.min(temp.vector)]]

print(result)
```

parLapply

```
library(snow)
data <- read.csv( 'dataset.csv' )
parallel.function <- function(i) {
  kmeans( data, centers=4, nstart=i )
}
cl <- makeCluster( mpi.universe.size(), type="MPI" )
clusterExport(cl, c('data'))

results <- parLapply( cl, c(25,25,25,25), fun=parallel.function )

temp.vector <- sapply( results, function(result) { result$tot.withinss } )
result <- results[[which.min(temp.vector)]]
print(result)

stopCluster(cl)
mpi.exit()
```

doSNOW foreach

```
library(foreach)
library(doSNOW)
data <- read.csv('dataset.csv')
cl <- makeCluster( mpi.universe.size(), type='MPI' )
clusterExport(cl,c('data'))
registerDoSNOW(cl)
results <- foreach( i = c(25,25,25,25) ) %dopar% {
  kmeans( x=data, centers=4, nstart=i )
}

temp.vector <- sapply( results, function(result)
  { result$tot.withinss } )
result <- results[[which.min(temp.vector)]]

print(result)
stopCluster(cl)
mpi.exit()
```

snow

```
cluster <- getMPIcluster()

# Print the hostname for each cluster member
sayhello <- function()
{
  info <- Sys.info()[c("nodename", "machine")]
  paste("Hello from", info[1], "with CPU type", info[2])
}

names <- clusterCall(cluster, sayhello)
print(unlist(names))

# stopCluster will call mpi.finalize, no need for mpi.exit
stopCluster(cluster)
```


Installing snow on ARCHER2

- Requires a slightly modified **Rmpi** build to install:
 - Firstly, we need to set up a modified ARCHER2 R build environment (configuring the compilers):


```
mkdir ~/.R
```

 - Add this to the file:


```
~/.R/Makevars
```
 - With the contents:


```
CC = cc
CXX = CC
FC = ftn
```
 - **Rmpi** install:
 - Make sure we have the R module loaded:


```
module load cray-R
```
 - Make sure we have the GNU compilers loaded:


```
module load PrgEnv-gnu
```
 - And we have the R library path setup to install on to the /work filesystem, i.e.:


```
export R_LIBS_USER=/work/ta055/ta055/$USER/Rinstall
```
 - Build **Rmpi**

```
R CMD INSTALL /work/z19/shared/adrianj/Rmpi_0.6-9.2.tar.gz --configure-args=" --with-Rmpi-type=CRAY"
```
- Finally, install snow:


```
R -e 'install.packages("snow", repos="https://cran.ma.imperial.ac.uk/")'
```

Running snow

```
#!/bin/bash
#SBATCH --job-name=snow
#SBATCH --nodes=2
#SBATCH --tasks-per-node=128
#SBATCH --time=0:5:0
#SBATCH --partition=standard
#SBATCH --qos=short
#SBATCH --account=z19
#SBATCH --hint=nomultithread

module load cray-R
export R_LIBS_USER=/work/z19/z19/adrianj/Rinstall
export PATH=$PATH:/work/z19/z19/adrianj/Rinstall/snow

export OMP_NUM_THREADS=1

srun RMPISNOW < ./simple_parallel_snow.R
```

More complex snow

```
nmax <- as.numeric(Sys.getenv("SLURM_NPROCS"))
cl <- getMPIcluster()
pbday <- function(n) {
  ntests <- 100000
  pop <- 1:365
  anydup <- function(i)
    any(duplicated(sample(pop, n, replace=TRUE)))
  sum(sapply(seq(ntests), anydup)) / ntests
}
clusterExport(cl, list('pbday'))

# print the time to do nmax tests, after distributing them to the workers
system.time( x <- clusterApply(cl, 1:nmax, function(n) { pbday(n) }) )

# compute the theoretical probability for each n
prob <- rep(0.0, nmax)
probnnot <- 1.0
for (i in 2:nmax) {
  probnnot <- probnnot*(366.0-i)/365.0
  prob[i] = 1.0 - probnnot
}

# print results, comparing tests to theory
z <- cbind(x, prob)
print(z)

# always include the following to stop the cluster
stopCluster(cl)
```

```
#!/bin/bash
#SBATCH --job-name=snow
#SBATCH --nodes=1
#SBATCH --tasks-per-node=64
#SBATCH --time=0:5:0
#SBATCH --partition=standard
#SBATCH --qos=short
#SBATCH --account=z19
#SBATCH --hint=nomultithread

module load cray-R
export R_LIBS_USER=/work/z19/z19/adrianj/Rinstall
export PATH=$PATH:/work/z19/z19/adrianj/Rinstall/snow

export OMP_NUM_THREADS=1

srun RMPISNOW < ./sample_birthday_snow.R
```

Summary

- R available and optimised on ARCHER2
- Controlling number and placement of workers can help with performance
- Many different mechanisms for parallel R and using multiple nodes
 - Can be configured to scale tasks, access more memory, reduce runtime
 - Choosing the most appropriate approach can improve performance